

Chap 4. Inter-Process Communication

■ Road map

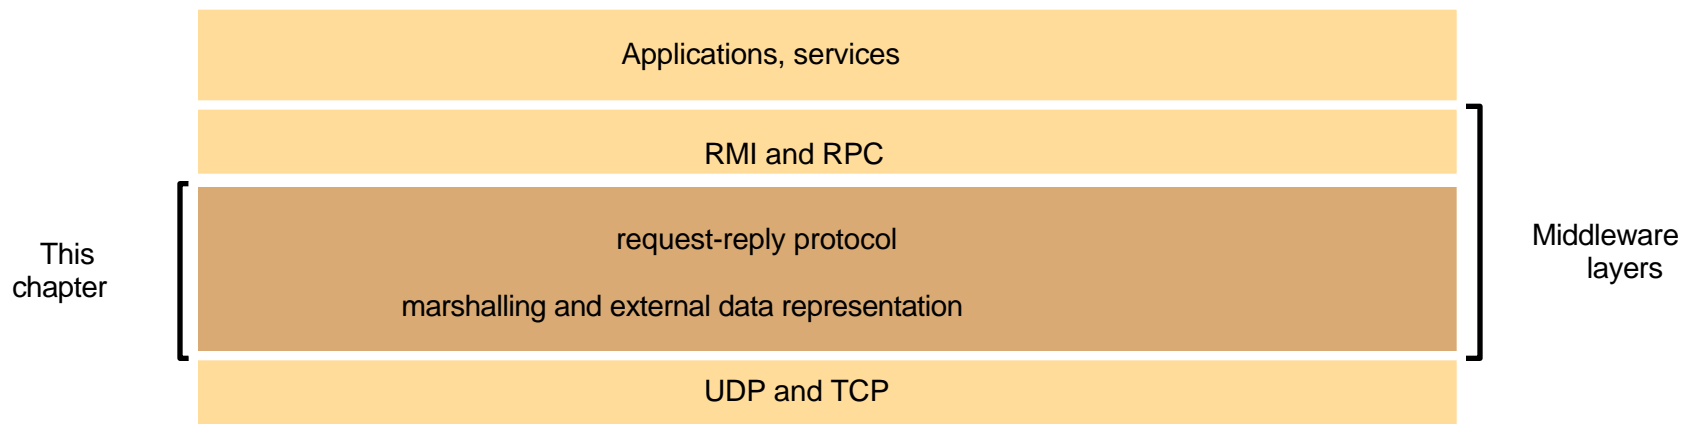
- ◆ 4.1. Intro
- ◆ 4.2. API for the Internet Protocols
- ◆ 4.3. External data representation and marshalling
- ◆ 4.4. Client-Server Communication
- ◆ 4.5. Group communication (self-read)
- ◆ 4.6. Case study (self-read)

4.1. Intro

- Focus:
 - ◆ Characteristics of protocols for communication between processes to model distributed computing architecture
 - ★ Effective means for communicating objects among processes at language level
 - ◆ Java API
 - ★ Provides both datagram and stream communication primitives/interfaces – building blocks for communication protocols
 - ◆ Representation of objects
 - ★ providing a common interface for object references
 - ◆ Protocol construction
 - ★ Two communication patterns for distributed programming: C-S using RMI/RPC and Group communication using 'broadcasting'

4.1. Intro

- In Chapter 3, we covered Internet transport (TCP/UDP) and network (IP) protocols without emphasizing how they are used at programming level
- In Chapter 5, we cover RMI facilities for accessing remote objects' methods AND the use of RPC for accessing the procedures in a remote server
- Chapter 4 is on how TCP and UDP are used in a program to effect communication via socket (e.g. Java sockets) – the Middle Layers – for object request/reply invocation and parameter marshalling/representation



4.2. API for internet protocols

■ IPC primitives

- ◆ message passing between a pair of processes can be supported by two message communication operations: *send* and *receive*
- ◆ *send* (destination, &msg); *receive* (source, &buf)
- ◆ Destination and source can be process id or port number (single receiver); Or mailbox (multiple receivers)
 - ★ Typically, (IP, port#) pair

- Use of sockets as API for UDP and TCP implementation – much more specification can be found at *java.net*

Message:

Header	Body
--------	------

4.2. API for internet protocols

- Synchronous communication
 - ◆ *Send* and *receive* processes synchronize at every message
 - ◆ Both blocking:
 - ★ whenever a *send* is issued, blocks until corresponding *receive* is issued; whenever a *receive* is issued, blocks until message arrives
- Asynchronous communication
 - ◆ *Send* from client is non-blocking, proceeds as soon as the message has been copied to a local buffer
 - ◆ *Receive* could be non-blocking or blocking, the latter has no disadvantages in system environment supporting multi-threading like Java
 - ★ **Non-blocking** receive is not very useful (you cannot proceed without message)

4.2. API for internet protocols

comparison

- Blocking

Advantages: Ease of use and low overhead of implementation

Disadvantage: low concurrency

- Non-blocking

Advantages: Flexibility, parallel operations


Disadvantages: Programming tricky: Program is timing-dependent (interrupt can occur at arbitrary time, and execution irreproducible)



→ Use blocking versions with **multiple threads**

4.2. API for internet protocols

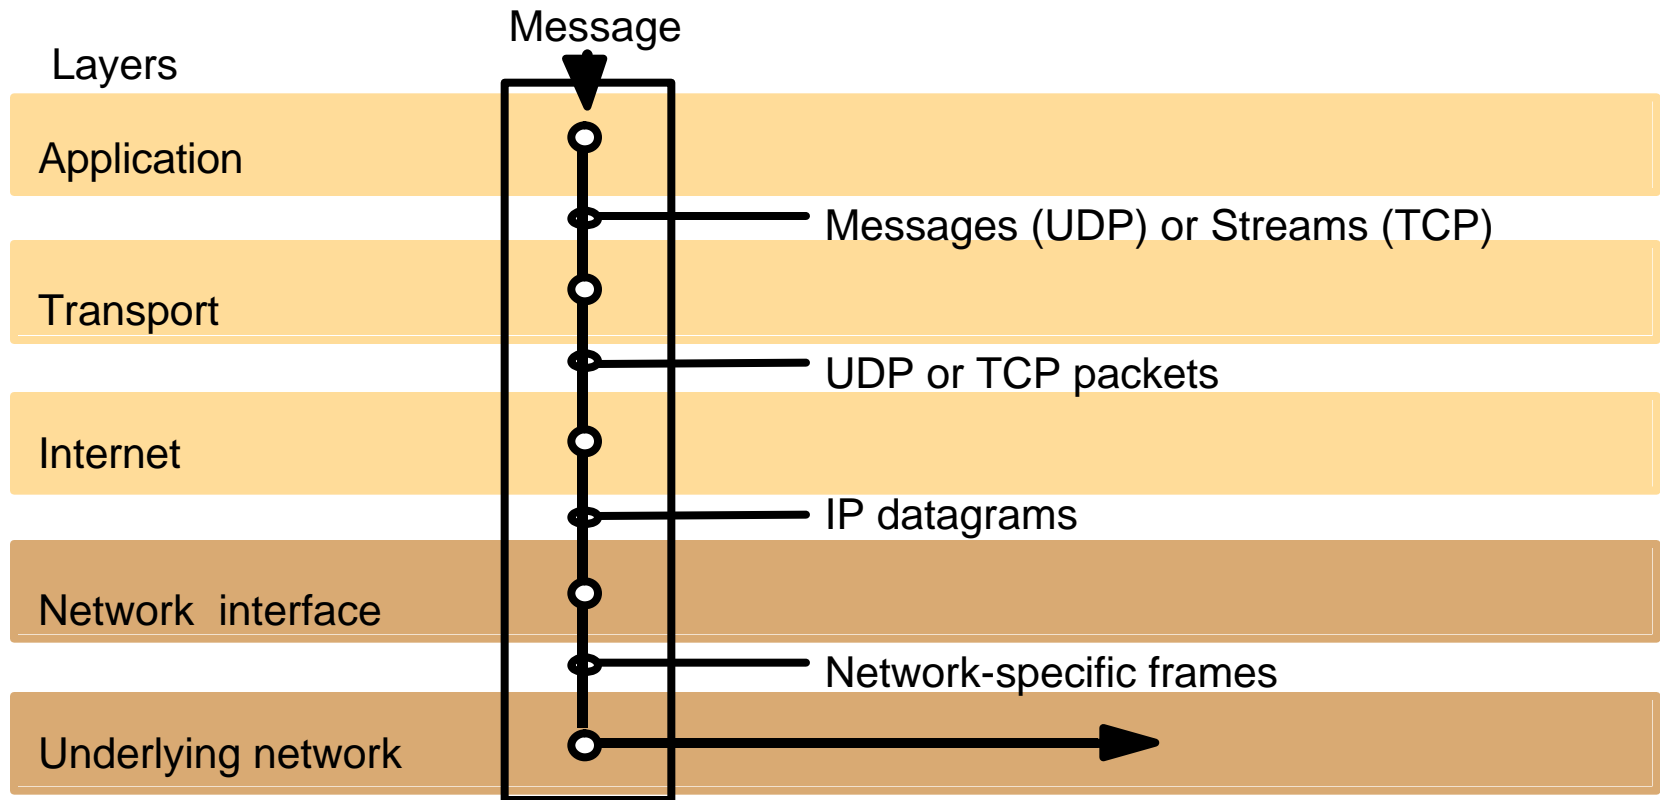
Using blocking operation without penalty



- Some threads may be blocked while others continue to be active


4.2. API for internet protocols

■ TCP/IP layers



4.2. API for internet protocols

- The programmer's conceptual view of a TCP/IP Internet



Recall: UDP, TCP, API

- UDP (User Datagram Protocol): offers no guarantee of delivery, a datagram protocol
- TCP (Transmission Control Protocol): reliable connection-oriented protocol, establishment of bi-directional communication channel
- API (Application programming interface)


4.2. API for internet protocols

■ Message destinations

- ◆ Messages are sent to a pair (Internet address, local port)
- ◆ Local port: an integer, message destination within a computer
 - ★ Each computer has 2^{16} possible ports available to local processes for receiving messages
 - ★ 0-1023: well-known for restricted use of privileged processes
 - ★ 1024-49151: registered ports, holds service descriptions
 - ★ 49152-65535: private purposes
 - ★ In practice, non-restricted ports can be used for private purposes, then, cannot use registered services simultaneously
- ◆ a port has exactly one receiver (except for multicast ports)
- ◆ receiving process can have many ports for different message types
- ◆ server processes usually publish their service-ports for clients

4.2. API for internet protocols

- **Sockets** : originated from BSD Unix
 - ◆ Provide an abstraction of endpoints for both TCP and UDP communication
 - ◆ Inter-process communication consists of transmitting a message between a socket in one process and a socket in another process
 - ◆ A socket must be bound to a local port and an IP address
 - ◆ Processes may use the same socket for sending and receiving messages
 - ◆ Sockets are typed/associated with a particular protocol, either TCP or UDP



Internet address = 138.37.94.248

Internet address = 138.37.88.249

4.2. API for internet protocols

- Java API for internet protocols
 - ◆ For either TCP or UDP, Java provides an `InetAddress` class, which contains a method: `getByName(DNS)` for obtaining IP addresses, irrespective of the number of address bits (32 bits for IPv4 or 128 bits for IPv6) by simply passing the DNS hostname. For example, a user Java code invokes:

```
InetAddress aComputer =  
InetAddress.getByName("www.cs.sfu.ca")
```

- ◆ The class encapsulates the details of the representation of the IP address

4.2. API for internet protocols

- UDP Datagram communication
 - ◆ Steps:
 - ★ Client finds an available port for UDP connection
 - ★ Client binds the port to local IP (obtained from `InetAddress.getByName(DNS)`)
 - ★ Server finds a designated port, publicizes it to clients, and binds it to local IP
 - ★ Server process issues a receive method and gets the IP and port # of sender (client) along with the message
 - ◆ Issues
 - ★ **Message size** – receiver needs to specify a buffer of certain size to receive a message. If message too big, truncated on arrival
 - ★ **Blocking** – send is non-blocking, returns when the message gets the UDP and IP layers; receive is blocking until a datagram is received or timeout
 - ★ **Timeouts** – reasonably large time interval can be set on receiver sockets to avoid indefinite blocking if required by program
 - ★ **Receive from any** – no specification of sources (senders)

4.2. API for internet protocols

- UDP Failure Models:
 - ◆ Omission failure: due to omission of send or receive (either checksum error or no buffer space at source or destination)
 - ◆ Ordering failure: due to out-of-order delivery
 - ◆ Applications using UDP are left to provide their own checks to achieve the quality of reliable communication they require (why and how?)
 - ★ UDP lacks built-in checks
 - ★ but failure can be modeled by implementing an ACK mechanism

4.2. API for internet protocols

UDP client sends a message to the server & gets a reply

```
import java.net.*; //defines socket-related classes
import java.io.*; //defines stream-related classes
public class UDPClient{
    public static void main(String args[]){ // args give message contents and server hostname
        DatagramSocket aSocket = null;
        try {
            aSocket = new DatagramSocket();
            byte [] m = args[0].getBytes();
            InetAddress aHost = InetAddress.getByName(args[1]);
            int serverPort = 6789;
            DatagramPacket request = new DatagramPacket(m, args[0].length(), aHost, serverPort);
            aSocket.send(request);
            byte[] buffer = new byte[1000];
            DatagramPacket reply = new DatagramPacket(buffer, buffer.length);
            aSocket.receive(reply);
            System.out.println("Reply: " + new String(reply.getData()));
        } catch (SocketException e) {System.out.println("Socket: " + e.getMessage());}
        } catch (IOException e) {System.out.println("IO: " + e.getMessage());}
    } finally { if(aSocket != null) aSocket.close();}
}
}
```

4.2. API for internet protocols

UDP server repeatedly receives a request and sends it back to the client

```
import java.net.*;
import java.io.*;
public class UDPServer{
    public static void main(String args[]){
        DatagramSocket aSocket = null;
        try{
            aSocket = new DatagramSocket(6789);
            byte[] buffer = new byte[1000];
            while(true){
                DatagramPacket request = new DatagramPacket(buffer, buffer.length);
                aSocket.receive(request);
                DatagramPacket reply = new DatagramPacket(request.getData(),
                    request.getLength(), request.getAddress(), request.getPort());
                aSocket.send(reply);
            }
        }catch (SocketException e){System.out.println("Socket: " + e.getMessage());}
        }catch (IOException e) {System.out.println("IO: " + e.getMessage());}
    }finally {if(aSocket != null) aSocket.close();}
}
```


4.2. API for internet protocols

- TCP Stream Communication issues:
 - ◆ Message sizes – user application has option to choose how much data it writes to a stream or reads from it
 - ◆ Lost messages / Flow control: dealt with by TCP
 - ◆ Message duplication and ordering – message identifiers are associated with each IP packet so as for recipient to detect and reject duplicates or re-order
 - ◆ Message destinations – a connection is established first. Once established, no IP addresses in packets needed (Each connection socket is bidirectional – using two streams: output/write and input/read)
 - ★ To establish a connection, client sends *connect* request to server, then server sends *accept* request to client
 - ★ Could be a overhead for a single C-S request and reply

4.2. API for internet protocols

■ TCP Stream Communication: other issues

- ◆ Matching of data items – both client/sender and server/receiver must agree on data types in the stream
- ◆ Threads – server creates a separate thread in accepting a connection, then it can block waiting for input without delaying other clients

■ Failure Model

- ★ Integrity: uses checksums for detection/rejection of corrupt data and seq #s for detection/rejection of duplicates
- ★ Validity: uses timeout with retransmission techniques to take care of packet losses
- ◆ Uses – TCP sockets used for such services as: HTTP, FTP, Telnet, SMTP

4.2. API for internet protocols

TCP client makes connection to server, sends request and receives reply

```
import java.net.*; //Defines socket-related classes
import java.io.*; //Defines stream-related classes
public class TCPCClient {
    public static void main (String args[]) {
        // arguments supply message and hostname of destination
        Socket s = null;
        try{
            int serverPort = 7896;
            s = new Socket(args[1], serverPort);
            DataInputStream in = new DataInputStream( s.getInputStream());
            DataOutputStream out = new DataOutputStream( s.getOutputStream());
            out.writeUTF(args[0]);          // UTF is a string encoding see Sn 4.3
            String data = in.readUTF();
            System.out.println("Received: "+ data) ;
        }catch (UnknownHostException e){
            System.out.println("Sock:"+e.getMessage());
        }catch (EOFException e){System.out.println("EOF:"+e.getMessage());}
        }catch (IOException e){System.out.println("IO:"+e.getMessage());}
    }finally {if(s!=null) try {s.close();}catch (IOException
e){System.out.println("close:"+e.getMessage());}}
    }
}
```

4.2. API for internet protocols

TCP server makes a connection for each client and then echoes the client's request

```
import java.net.*;
import java.io.*;
public class TCPServer {
    public static void main (String args[]) {
        try{
            int serverPort = 7896;
            ServerSocket listenSocket = new ServerSocket(serverPort);
            while(true) {
                Socket clientSocket = listenSocket.accept();
                Connection c = new Connection(clientSocket);
            }
        } catch(IOException e) {System.out.println("Listen :"+e.getMessage());}
    }
}

// this figure continues on the next slide
```

4.2. API for internet protocols

```
class Connection extends Thread {
    DataInputStream in;
    DataOutputStream out;
    Socket clientSocket;
    public Connection (Socket aClientSocket) {
        try {
            clientSocket = aClientSocket;
            in = new DataInputStream( clientSocket.getInputStream());
            out =new DataOutputStream( clientSocket.getOutputStream());
            this.start();
        } catch(IOException e) {System.out.println("Connection:"+e.getMessage());}
    }
    public void run(){
        try {
            // an echo server
            String data = in.readUTF();
            out.writeUTF(data);
        } catch(EOFException e) {System.out.println("EOF:"+e.getMessage());}
        } catch(IOException e) {System.out.println("IO:"+e.getMessage());}
        } finally{ try {clientSocket.close();}catch (IOException e){/*close failed*/}}
    }
}
```